

Graduation Project - Total Variational Image Denoising with U-Net

Thanh Trung Vu - 230849442

1 Introduction

In this thesis we will describe the theory and application of the total variational (TV) method to image denoising. We will also describe the implementation of the method and an improvement using a U-Net architecture to find the regularisation parameters more efficiently.

Image processing, and image denoising in particular, has seen significant advancements since the U-Net was first applied in 2015. This project aims to build upon these advancements by implementing and experimenting with a U-Net-based model for image denoising. In particular, it combines U-Net with a traditional image-denoising algorithm called the Primal Dual Hybrid Gradient (PDHG) algorithm [1]. Specifically, a PDHG solver is appended to the end of the U-Net, treating the solver as the final layer of the entire network to create an end-to-end unsupervised model.

One notable aspect of using U-Net in this problem is that, instead of acting as a black box and outputting a denoised version of each input image, it outputs a parameter map used in another algorithm, demonstrating the flexibility of neural networks.

1.1 Denoising problems

Total variational (TV) method is one of the most popular methods for image denoising. The TV method works by finding the image that has the smallest total variation and is closest to the noisy image in terms of the mean squared error.

The TV method is based on the idea that images that are smooth have a small total variation, while images that are noisy have a large total variation. By minimising the total variation of the image, the TV method is able to remove the noise from the image and produce a smoother image [4].

Here we combine the traditional TV method with the U-Net to produce a more efficient method for image denoising.

The problem with TV method is that it is not clear how to choose the regularisation parameter λ . The U-Net can be used to find the regularisation parameter λ more efficiently. After training with many examples, the U-Net can learn to predict the regularisation parameter λ that produces the best denoised image for any particular noisy image.

1.2 Total Variational Denoising

The total variational (TV) method is a method for image denoising that is based on the principle of minimising the total variation of the image. The total variation of an image is a measure of the amount of variation in the intensity of the image. The TV method works by finding the image that has the smallest total variation and is closest to the noisy image in terms of the mean squared error.

The TV method is based on the idea that images that are smooth have a small total variation, while images that are noisy have a large total variation. By minimising the total variation of the image, the TV method is able to remove the noise from the image and produce a smoother image.

Denoising problems can be described as

$$\mathbf{z} = \mathbf{x}_{\text{true}} + \mathbf{e} \quad (1)$$

where $\mathbf{x}_{\text{true}} \in \mathbb{R}^{m \times n}$ is the object to be imaged, $\mathbf{e} \in \mathbb{R}^{m \times n}$ denotes some random noise component, and $\mathbf{z} \in \mathbb{R}^{m \times n}$ represents the measured data. The goal is to recover \mathbf{x}_{true} or at least a good enough reconstruction given \mathbf{z} .

A prominent approach is to formulate the reconstruction as a minimisation problem:

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} d(\mathbf{x}, \mathbf{z}) + \mathcal{R}(\mathbf{x}) \quad (2)$$

where $d(\cdot, \cdot)$ is a data fidelity term and $\mathcal{R}(\cdot)$ is a regularisation term.

Here we will look at Gaussian noise. The data fidelity term appropriate for Gaussian noise is the square of the ℓ_2 norm:

$$d(\mathbf{x}, \mathbf{z}) = \frac{1}{2} \|\mathbf{x} - \mathbf{z}\|_2^2 \quad (3)$$

The regularisation term $\mathcal{R}(\cdot)$ is the total variation of the image, which is a measure of the amount of variation in the intensity of the image. The total variation of an image can be calculated using the following formula:

$$\mathcal{R}(\mathbf{x}) = \lambda \|\nabla \mathbf{x}\|_1 \quad (4)$$

where $\nabla \mathbf{x}$ is the gradient of the image \mathbf{x} and $\lambda \in \mathbb{R}^+$ is a regularisation parameter that controls the trade-off between the total variation of the image and the mean squared error between the denoised image \mathbf{x} and the noisy image \mathbf{z} .

∇ denotes a finite-differences operator. Finite difference operators include, forward difference operator, backward difference operator, shift operator, central difference operator and mean operator. Let us focus on the forward difference operator. In particular, in the discrete case, given a matrix \mathbf{x} , $\nabla \mathbf{x}$ is two matrices, one for the vertical gradient $\nabla_x \mathbf{x}$ and one for the horizontal gradient $\nabla_y \mathbf{x}$.

$$\nabla \mathbf{x} = \begin{bmatrix} \nabla_x \mathbf{x} \\ \nabla_y \mathbf{x} \end{bmatrix} \quad (5)$$

For the forward gradient, the vertical gradient is calculated as:

$$\begin{aligned}
\nabla_x \mathbf{x} &= \begin{bmatrix} \mathbf{x}_{2,1} - \mathbf{x}_{1,1} & \mathbf{x}_{2,2} - \mathbf{x}_{1,2} & \dots & \mathbf{x}_{2,n-1} - \mathbf{x}_{1,n-1} & \mathbf{x}_{2,n} - \mathbf{x}_{1,n} \\ \mathbf{x}_{3,1} - \mathbf{x}_{2,1} & \mathbf{x}_{3,2} - \mathbf{x}_{2,2} & \dots & \mathbf{x}_{3,n-1} - \mathbf{x}_{2,n-1} & \mathbf{x}_{3,n} - \mathbf{x}_{2,n} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ \mathbf{x}_{m,1} - \mathbf{x}_{m-1,1} & \mathbf{x}_{m,2} - \mathbf{x}_{m-1,2} & \dots & \mathbf{x}_{m,n-1} - \mathbf{x}_{m-1,n-1} & \mathbf{x}_{m,n} - \mathbf{x}_{m-1,n} \\ \mathbf{x}_{1,1} - \mathbf{x}_{m,1} & \mathbf{x}_{1,2} - \mathbf{x}_{m,2} & \dots & \mathbf{x}_{1,n-1} - \mathbf{x}_{m,n-1} & \mathbf{x}_{1,n} - \mathbf{x}_{m,n} \end{bmatrix} \\
&= \begin{bmatrix} -1 & 1 & \dots & 0 & 0 \\ 0 & -1 & \dots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & -1 & 1 \\ 1 & 0 & \dots & 0 & -1 \end{bmatrix} \mathbf{x}
\end{aligned} \tag{6}$$

and the horizontal gradient is calculated as:

$$\begin{aligned}
\nabla_y \mathbf{x} &= \begin{bmatrix} \mathbf{x}_{1,2} - \mathbf{x}_{1,1} & \mathbf{x}_{1,3} - \mathbf{x}_{1,2} & \dots & \mathbf{x}_{1,n} - \mathbf{x}_{1,n-1} & \mathbf{x}_{1,1} - \mathbf{x}_{1,n} \\ \mathbf{x}_{2,2} - \mathbf{x}_{2,1} & \mathbf{x}_{2,3} - \mathbf{x}_{2,2} & \dots & \mathbf{x}_{2,n} - \mathbf{x}_{2,n-1} & \mathbf{x}_{2,1} - \mathbf{x}_{2,n} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ \mathbf{x}_{m-1,2} - \mathbf{x}_{m-1,1} & \mathbf{x}_{m-1,3} - \mathbf{x}_{m-1,2} & \dots & \mathbf{x}_{m-1,n} - \mathbf{x}_{m-1,n-1} & \mathbf{x}_{m-1,1} - \mathbf{x}_{m-1,n} \\ \mathbf{x}_{m,2} - \mathbf{x}_{m,1} & \mathbf{x}_{m,3} - \mathbf{x}_{m,2} & \dots & \mathbf{x}_{m,n} - \mathbf{x}_{m,n-1} & \mathbf{x}_{m,1} - \mathbf{x}_{m,n} \end{bmatrix} \\
&= \mathbf{x} \begin{bmatrix} -1 & 0 & \dots & 0 & 1 \\ 1 & -1 & \dots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & -1 & 0 \\ 0 & 0 & \dots & 1 & -1 \end{bmatrix}
\end{aligned} \tag{7}$$

We calculate the gradient in the horizontal and vertical directions separately, and then combine them to get the total variation of the image.

$$\|\nabla \mathbf{x}\|_1 = \|\nabla_x \mathbf{x}\|_1 + \|\nabla_y \mathbf{x}\|_1 \tag{8}$$

For the forward gradient, the norm-1 of the vertical gradient is calculated as:

$$\|\nabla_x \mathbf{x}\|_1 = \left(\sum_{i=1}^{m-1} \sum_{j=1}^n |\mathbf{x}_{i+1,j} - \mathbf{x}_{i,j}| \right) + \sum_{j=1}^n |\mathbf{x}_{1,j} - \mathbf{x}_{m,j}| \tag{9}$$

and norm-1 of the horizontal gradient is calculated as:

$$\|\nabla_y \mathbf{x}\|_1 = \left(\sum_{i=1}^m \sum_{j=1}^{n-1} |\mathbf{x}_{i,j+1} - \mathbf{x}_{i,j}| \right) + \sum_{i=1}^m |\mathbf{x}_{i,1} - \mathbf{x}_{i,n}| \tag{10}$$

where $\mathbf{x}_{i,j}$ is the intensity value of pixel (i, j) in the image \mathbf{x} . Note that $\mathbf{x}_{i,j}$ can be a scalar for a black-and-white image or a vector of 3 values for a colour image.

Note that, we can move λ inside the norm, so the regularisation term can be written as:

$$\mathcal{R}(\mathbf{x}) = \lambda \|\nabla \mathbf{x}\|_1 = \|\lambda \nabla \mathbf{x}\|_1 = \|\lambda \nabla_x \mathbf{x}\|_1 + \|\lambda \nabla_y \mathbf{x}\|_1 \quad (11)$$

We can see that the regularisation parameter does not need to be a scalar.

$$\mathcal{R}(\mathbf{x}) = \|\mathbf{\Lambda} \circ \nabla \mathbf{x}\|_1 = \|\mathbf{\Lambda}^{(x)} \circ \nabla_x \mathbf{x}\|_1 + \|\mathbf{\Lambda}^{(y)} \circ \nabla_y \mathbf{x}\|_1 \quad (12)$$

where $\mathbf{\Lambda}^{(x)}$ and $\mathbf{\Lambda}^{(y)}$ are matrices that contain the regularisation parameters for the vertical and horizontal gradients respectively, and \circ denotes the Hadamard or element-wise product. We will refer to $\mathbf{\Lambda}$ as the regularisation parameter map.

Putting it all together, the solution to the total variational denoising problem can be defined using the following formula:

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \frac{1}{2} \|\mathbf{x} - \mathbf{z}\|_2^2 + \lambda \|\nabla \mathbf{x}\|_1 \quad (13)$$

if we use a scalar λ regularisation parameter, or

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \frac{1}{2} \|\mathbf{x} - \mathbf{z}\|_2^2 + \|\mathbf{\Lambda} \nabla \mathbf{x}\|_1 \quad (14)$$

if we use a regularisation parameter map $\mathbf{\Lambda}$ of regularisation parameters.

In this thesis, we will use U-Net to help us find the regularisation parameter map $\mathbf{\Lambda}$ given any given (noisy) image. We will compare the results to the best-case scenario where we use a scalar regularisation parameter λ .

1.3 U-Net

The U-Net is a type of neural network that is commonly used for image segmentation [3]. The U-Net is an encoder-decoder network that is designed to take an image as input and produce a segmentation mask as output. The U-Net is made up of a series of convolutional layers that downsample the image and a series of transposed convolutional layers that upsample the image. The U-Net is able to learn to segment images by training on a large dataset of images and their corresponding segmentation masks.

In our case, we will optimise a U-Net model to find the regularisation parameter λ that produces the best denoised image for any particular noisy image.

1.4 Related Work

1.5 Contribution

2 Task Description/Method/Implementation

2.1 Primal Dual Hybrid Gradient (PDHG)

To solve equation (14) we can use an iterative method called the Primal Dual Hybrid Gradient (PDHG).

The algorithm is

Algorithm 1 PDHG algorithm for image denoising with fixed regularization parameter-map Λ (adapted from ...)

```

1: Input:  $L = \|\nabla\|^T$ ,  $\tau = 1/L$ ,  $\sigma = 1/L$ ,  $\theta = 1$ , noisy image  $\mathbf{x}_0$ 
2: Output: reconstructed image  $\hat{\mathbf{x}}$ 
3:  $\bar{\mathbf{x}}_0 = \mathbf{x}_0$ 
4:  $\mathbf{p}_0 = 0$ 
5:  $\mathbf{q}_0 = 0$ 
6: for  $k < T$  do
7:    $\mathbf{p}_{k+1} = (\mathbf{p}_k + \sigma(\bar{\mathbf{x}}_k - \mathbf{x}_0)) / (1 + \sigma)$ 
8:    $\mathbf{q}_{k+1} = \text{clip}_{\Lambda}(\mathbf{q}_k + \sigma \nabla \bar{\mathbf{x}}_k)$ 
9:    $\mathbf{x}_{k+1} = \mathbf{x}_k - \tau \mathbf{p}_{k+1} - \tau \nabla^T \mathbf{q}_{k+1}$ 
10:   $\bar{\mathbf{x}}_{k+1} = \mathbf{x}_{k+1} + \theta(\mathbf{x}_{k+1} - \mathbf{x}_k)$ 
11: end for
12:  $\hat{\mathbf{x}} = \mathbf{x}_T$ 

```

where clip_{Λ} is a function that clips the values of \mathbf{q}_{k+1} to the range $[0, \Lambda]$.

The challenge lies in finding or learning the regularisation parameter map Λ . Using a single λ value penalises all regions of the image equally, which is suboptimal. Instead, it would be better to assign higher penalties to smooth regions and lower penalties to detailed regions to preserve image details. However, this approach is not perfect, as it might leave noise in detailed areas.

Since regularisation is multiplied by the gradient, higher λ values will penalize areas with many changes (typically areas with more details), while lower λ values prioritize detail preservation over noise reduction.

The performance of Primal-Dual Hybrid Gradient (PDHG) methods can be improved by assigning suitable regularisation parameters to regions based on their characteristics. However, narrowing down the search space is challenging, as we potentially have $d^{(n^2)}$ different maps to consider, where d is the number of λ values (typically around 100, ranging from 0.01 to 1) and n is the image dimension. This is computationally infeasible.

To address this, we can apply image segmentation. Effective image segmentation helps determine appropriate λ values for different regions. A good network for learning image segmentation is U-Net [3].

3 Architecture

For this project, our model comprises two primary components:

1. A U-Net, denoted as NET_{Θ} , which is responsible for learning the regularisation parameters Λ_{Θ} from an input image \mathbf{x}_0 .
2. A Primal Dual Hybrid Gradient (PDHG) algorithm solver (refer to Appendix), which utilizes both \mathbf{z} and Λ_{Θ} to transform \mathbf{x}_0 into $\hat{\mathbf{x}}$. This transformation aims to minimise the loss function $L(\mathbf{z}, \mathbf{x})$, thereby facilitating effective reconstruction. The symbol Θ represents the set of trainable parameters within the network.




Figure 1: The general architecture

We interpret the PDHG solver as the ultimate hidden layer within our network. The outcome of the loss function $L(\mathbf{z}, \mathbf{x})$ is subsequently employed for backpropagation, essential for training the parameters Θ .

For our implementation, we consider $n \times n$ images, specifically with $n = 512$.

In terms of variables:

- $\mathbf{z} \in \mathbb{R}^{n^2}$ represents the clean image, serving as the ground truth.
- $\mathbf{x}_0 \in \mathbb{R}^{n^2}$ denotes the noisy image inputted to NET_Θ .
- Λ_Θ , the regularisation parameters map, emerges as the final output from the U-Net and is instrumental for the PDHG algorithm solver. Throughout this project, Λ and Λ_Θ are used interchangeably.

The architecture of NET_Θ is inspired by the U-Net, a specialised type of convolutional neural network architecture.

Let us first delve into the foundational concept of a convolutional neural network, which underpins the U-Net architecture.

3.1 Convolutional Neural Network (CNN)

Convolutional Neural Networks (CNNs) are specialised types of artificial neural networks optimised for analyzing image data and other grid-like structures. We base our understanding on a two-dimensional (2D) CNN architecture, where each layer is comprised of one or more 2D matrices, in contrast to the vectors used in standard "vanilla" neural networks.

Inspiration and Function

The design of CNNs draws inspiration from the human visual cortex, which processes visual stimuli by sequentially extracting features, ranging from simple edges and shapes to complex objects. CNNs replicate this hierarchical feature extraction using convolutional layers and pooling layers.

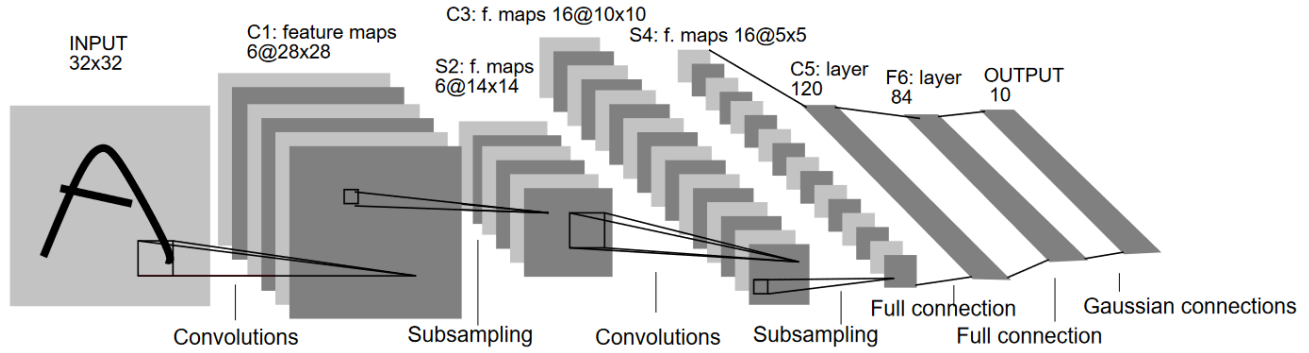


Figure 2: Architecture of LeNet-5, one of the earliest Convolutional Neural Networks [2]

Convolutional Layers

In a convolutional layer, each output element results from a linear combination of inputs, processed through an activation function—akin to "vanilla" neural networks. However, the critical difference lies in the utilization of filters (or kernels). Each filter traverses the input matrix, computes the dot product with the corresponding input elements, and channels the result through an activation function. A single filter generates one output matrix, termed a feature map, acting as a feature extractor that emphasizes specific characteristics of the input. When multiple filters are employed in a layer, they produce multiple feature maps, with each map corresponding to a different channel. This design significantly reduces the proportion of trainable parameters relative to the number of output values, enhancing efficiency for image-related tasks.

Pooling Layers

An essential component of CNNs is the pooling layer, typically implemented as a max-pooling layer. This layer simplifies the output by retaining only the maximum values from the defined regions in the feature map. The objective is to preserve the most significant features while reducing data dimensionality. In this project, max-pooling is a crucial technique used to enhance model performance.

The convolutional and pooling layers constitute the foundational blocks of the U-Net architecture.

3.2 U-Net

A U-Net, as introduced by Ronneberger et al. [3], extends the conventional CNN architecture by incorporating skip connections, akin to those in a residual network. These skip connections concatenate the outputs of distinct convolutional layers to generate a larger composite output.

Encoder

The U-Net architecture is divided into two principal components: an Encoder and a Decoder. The Encoder functions similarly to a standard CNN, utilizing a combination of convolutional layers and max-pooling layers organized into distinct "blocks." In this project, each Encoder block comprises a pair of convolutional layers followed by a max-pooling layer, designed to successively double the number of channels (or feature maps) while concurrently reducing the feature map size due to the pooling.

Decoder

Conversely, the Decoder is structured around a series of blocks that progressively halve the number of channels while enlarging the feature map size, effectively "unrolling" the data. Unlike the Encoder blocks, each Decoder block concludes with an up-convolutional layer followed by a skip connection. This up-convolutional layer, essentially a regular convolutional layer, outputs data that is then concatenated with the output from a corresponding Encoder block, facilitating the feature integration across different layers of the network.




Figure 3: Our NET_θ Architecture

In our model, the Encoder consists of three encoding blocks. First we increase the number of channels to 32. The first block the double the number of channels to 64 feature maps. The second and third blocks increase the number of channels further to 128 and 256, respectively.

Similarly we have three decoding blocks. Each decoding block in our model concludes with an up-convolutional layer, which serves the same function as a typical convolutional layer, followed by a skip connection. The first skip connection takes the 32-channel output from the second encoding block and merges it with the 32-channel output from the first up-convolutional layer. This combination produces a 64-channel output, which is then processed by the first decoding block. The second skip connection similarly merges the 16-channel output from the first encoding block with the 16-channel output from the second up-convolutional layer, resulting in a 32-channel input for the second decoding block.

The final stage in our architecture employs a convolutional layer to convert the 16-channel output from the last decoding block into a 2-channel output, thus producing two matrices as the ultimate output of our model.

4 Experimental Set-Up

We used the Chest X-Ray Images (Pneumonia) dataset which was downloaded from <https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia>. This is a dataset for binary classification of normal and pneumonia chest X-ray images. The dataset consists of 5,863 X-Ray images (JPEG) and 2 categories (Pneumonia/Normal). All images are black-and-white and are in the JPEG format. The X-rays images have various resolutions.

We used 800 images for the training dataset, 100 images for the validation dataset, and 100 images for the test dataset. We cropped the images to squares and then resized all images to 256x256 pixels. For each image we generated a random value of σ ranging from 0.1 to 0.5 and noised the image with Gaussian noise with that σ .

For learning the regularisation parameter map, we used a U-Net structure with 3 blocks and 32 initial filters to denoise the images, resulting in 1,796,034 trainable parameters. We used the Adam optimiser with an initial learning rate of $1e-4$. The batch size is set to 1. We used the Mean Square Error (MSE) loss function to train the U-Net.

We set T_{train} to 256 and T_{test} to 256. We trained for 100 epochs.

(The log can be found in "true-voice-32" https://wandb.ai/wof/chest_xray/runs/wedf9ud9)

(I also did an experiment where I removed PDHG completely and used only the UNet to learn how to denoise an image (https://wandb.ai/wof/chest_xray/runs/xnjgywx7/overview).)

5 Results

Each test image was also noised with Gaussian noise with a random value of σ ranging from 0.1 to 0.5.

We compare the results using the regularisation parameter map generated by our method with the results using a scalar λ for the total variational denoising method. For each test image, we run a grid search to find the best scalar λ . An example result is shown in Figure ???. For this particular test image, the best scalar λ was 0.02, which resulted in a PSNR of 35.99 dB. The regularisation parameter map generated by our method resulted in a PSNR of 37.26 dB. The resulting denoised images are shown in Figure ??.




Figure 4: Chest X-ray example - grid search




Figure 5: Chest X-ray example: clean, noisy with $\sigma = 0.5$, $\lambda = 0.07$ with PSNR = 29.945, Λ_θ with PSNR = **31.248**

	PDHG - λ	PDHG - Λ_Θ
SSIM	0.87 ± 0.04	0.90 ± 0.03
PSNR	33.43 ± 2.16	34.73 ± 2.02

Table 1: Mean and standard deviation of the measures SSIM and PSNR obtained over the test set for the chest x-ray dataset example. The TV-reconstruction using the parameter-map Λ_Θ improves both measures.




Figure 8: Histograms of the PSNR and SSIM differences between using λ scalar and using Λ_{Θ} map for reconstructions of noisy versions of the test images. The Λ_{Θ} map method consistently outperforms the scalar λ method.

Plotting the values for all the test images, we can see that the regularisation parameter map consistently outperforms the scalar λ . The box plots in Figure ?? show the distribution of PSNR and SSIM values for the test images using the two methods. The regularisation parameter map consistently produces higher PSNR and SSIM values than the scalar λ .

- (More experiments) Test with the sigma 0.1 and 0.3 for test images and add: box plots, table rows, and difference histograms.

6 Conclusion

This project showcases the integration of neural networks with traditional algorithms to enhance image denoising capabilities, marking a significant step forward in computational imaging.

6.1 Experimentation and Findings

We experimented with the capabilities of the U-Net architecture to learn and apply regularisation parameters dynamically across different image segments using the PDHG algorithm. Traditional denoising techniques typically rely on a single lambda value for regularisation, which often does not cater to the varying needs across an image. Our approach attempts to transcend this limitation by approximating an optimal lambda map through the U-Net, enhancing the PDHG’s performance. This methodology, while computationally intensive, illustrates the potential benefits of segment-specific regularisation in image denoising.

In conclusion, while the experiment demonstrated potential improvements over traditional single-lambda denoising methods, the current computational demands make it impractical for real-world applications.






7 Future Work

References

- [1] Andreas Kofler et al. *Learning Regularization Parameter-Maps for Variational Image Reconstruction using Deep Neural Networks and Algorithm Unrolling*. 2023. arXiv: [2301.05888](https://arxiv.org/abs/2301.05888) [math.OC]. URL: <https://github.com/koflera/LearningRegularizationParameterMaps>.
- [2] Y. Lecun et al. “Gradient-based learning applied to document recognition”. In: *Proceedings of the IEEE* 86.11 (1998), pp. 2278–2324. DOI: [10.1109/5.726791](https://doi.org/10.1109/5.726791). URL: https://www.researchgate.net/publication/2985446_Gradient-Based_Learning_Applied_to_Document_Recognition.
- [3] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. *U-Net: Convolutional Networks for Biomedical Image Segmentation*. 2015. arXiv: [1505.04597](https://arxiv.org/abs/1505.04597) [cs.CV].
- [4] Leonid I. Rudin, Stanley Osher, and Emad Fatemi. “Nonlinear total variation based noise removal algorithms”. In: *Physica D: Nonlinear Phenomena* 60.1 (1992), pp. 259–268. ISSN: 0167-2789. DOI: [https://doi.org/10.1016/0167-2789\(92\)90242-F](https://doi.org/10.1016/0167-2789(92)90242-F). URL: <https://www.sciencedirect.com/science/article/pii/016727899290242F>.

The following list contain the primary Python libraries used in the implementation of this project.

- **NumPy**: Utilised for numerical representations and operations, including vector and matrix multiplication.
- **PyTorch**: Employed for constructing and training neural network models due to its robust, flexible, and efficient computational dynamics.
- **matplotlib** and **seaborn**: Employed for data visualisation tasks, including the creation of histograms, box plots, and other graphical representations.
- **scikit-learn**: Used for benchmarking our model against traditional machine learning algorithms, providing tools for data preprocessing, cross-validation, and more.
- **random**: Critical for generating random numbers, used extensively in stochastic processes like K-fold cross-validation and bootstrap sampling.
- **wandb**: Integrated for experiment tracking and logging, facilitating the monitoring of model training metrics and performance in real-time.



(a) Λ_Θ map and histogram of the values

(b) $\log(\Lambda_\Theta + 0.05)$ and histogram of the log values

Figure 6: The produced Λ_Θ map for the chest x-ray dataset example. Since the values are too close to one another in the lower spectrum, the bottom row shows the log of the map and the histogram of the log values for better separation.




Figure 7: Test results summary